

Technical Roadmap: Path to State-of-the-Art (SOTA)

Executive Summary

The Problem We Solve

Cloud cost anomalies bleed 15-30% of your infrastructure budget. Native tools (AWS Cost Explorer, GCP Recommender) detect anomalies *after* they hit your bill—hours to days later. By then, the damage is done.

We detect anomalies *before* they become costs—on operational signals (API calls, error rates, latency) rather than billing data. This gives you hours of lead time to prevent runaway spend.


What We've Built (Current State)

Metric	Value	What It Means
PR-AUC	0.438	93x better than guessing (random = 0.0047)
Inference Speed	0.8 μ s	1.2M predictions/second on single core
False Positive Rate	Controlled	Won't spam on-call engineers

The Business Case

Client Scale	Annual Cloud Spend	Expected Savings	Your Investment
Startup	\$5–10M	\$500K–1M/yr (10–20%)	2-week integration
Mid-Market	\$10–50M	\$1–5M/yr	Replaces 1-2 FTE cloud engineers
Enterprise	\$100M+	\$10–20M+	SOX-compliant audit trails

Quick Wins vs. Big Bets



Funding Scenarios

Scenario	Investment	What You Get	Timeline	Expected Result
Minimal	2 weeks, 1 engineer	Quick wins (Phases 1-2)	2 weeks	PR-AUC: 0.52-0.56
Moderate	6 weeks, 1-2 engineers	Core improvements (Phases 1-4)	6 weeks	PR-AUC: 0.60-0.68
Full	4 months, 2-3 engineers	Complete system (Phases 1-7)	4 months	PR-AUC: 0.85+

The Ask

Start with Phases 1-3 (2 weeks, 1 engineer). Low risk, high return. If results validate, proceed to Phase 4+.

PART 1: CONTEXT & POSITIONING

Executive Summary: Performance Trajectory

Milestone	Technique	Expected PR-AUC	Timeline
Current	XGBoost + 51 features	0.438	Done
Phase 1	LightGBM + Cyclical + EMA	0.48–0.52	1–2 weeks
Phase 2	SMOTE + Focal Loss	0.52–0.56	1 week
Phase 3	Graph Features (HHI, Error Prop)	0.55–0.60	2 weeks
Phase 4	LSTM/Transformer	0.60–0.68	3–4 weeks
Phase 5	Ensemble + Conformal	0.65–0.72	2 weeks
Phase 6	Adaptive Baselines + Active Learning	0.70–0.78	2–3 weeks
Phase 7	Dual-Brain Architecture	0.85+ (Hybrid)	Month 4+
SOTA Target	Full Pipeline	0.75+	12–16 weeks
Theoretical Ceiling	Causal + Oracle Labels	0.85–0.95	Research frontier

Note: Gains are not strictly additive; diminishing returns expected. Theoretical limits remain beyond current practical systems.

Industry Benchmarks: PR-AUC in Context

Domain	Typical PR-AUC	Positive Rate	This Roadmap Target	Source
Credit Card Fraud	0.50–0.65	0.17%	0.75 (Phase 6)	Kaggle IEEE-CIS [10]
Network Intrusion	0.55–0.70	1–5%	0.75	Schmidl et al. 2022 [4]
Cloud Cost (AWS/GCP 2025)	0.60–0.70	~1%	0.85 (Phase 7)	AWS GA 2025 [7]
Behavioral Anomaly	0.40–0.60	0.5%	0.85+	This system

Current Performance Context:

- With **0.47% anomaly prevalence**, random baseline PR-AUC = **0.0047**
- Current PR-AUC of **0.438** = **93x lift over random** (production-grade)
- Phase 7 target of **0.85** = **181x lift** (research frontier)

How Native Cloud Cost Tools Work (Accurate Comparison)

Important Clarification: Native cloud cost anomaly tools are not "slow" in their ML inference—they are constrained by **billing data pipelines**.

AWS Cost Anomaly Detection

- Uses near-real-time *spend telemetry*, but underlying **billing data settles with delay** (hours to next day)
- Anomalies are evaluated on **finalized cost data** that can lag due to credits, refunds, and usage reconciliation [7]
- In practice: alerts typically arrive **same-day or next-day**, not always 24-48 hours, but not sub-second either

Google Cloud Cost Anomaly Detection

- Analyzes **daily and intra-day spend trends** tied to billing pipelines [6]
- Detection cadence is **hourly to daily**, not millisecond-level

Datadog / Third-Party Tools

- For **metrics/log anomalies** (APM, infra): can flag anomalies in **near real-time (seconds–minutes)** on streaming telemetry
- For **cost dashboards**: constrained by cloud provider billing data cadence (hour-scale lags)

Our Architectural Difference

Dimension	Native Cloud Cost Tools	This System
Data Source	Billing/aggregation pipelines	Operational telemetry (requests, errors, usage)

Dimension	Native Cloud Cost Tools	This System
Detection Latency	Hourly to daily (billing-constrained)	Sub-millisecond capable (0.8 μ s inference)
When Alert Fires	After spike appears in billing data	Before spike hits the bill
Trade-off	Accurate cost attribution	Early warning on leading indicators

The key insight: we operate on **operational signals** (upstream), not **billing signals** (downstream). This enables detection before cost data even settles.

Deployment Model: Detection-Only with Human-in-the-Loop

Critical Scope Clarification: This system is designed for **detection and alerting only**. We do not automatically modify customer infrastructure.

What We Provide

- Real-time anomaly **detection** and **scoring**
- Actionable **alerts** with root-cause explanations
- **Recommendations** for remediation (not automatic execution)

What We Do NOT Do

- Automatically pause/terminate customer resources
- Directly modify cloud configurations
- Execute remediation without human approval

Customer Integration Options

Option	Description	Use Case
Alert-Only	Anomaly alerts sent via Slack/PagerDuty/email	Conservative teams, production-critical systems
Human-in-the-Loop	Alert + one-click remediation approval	Balanced speed/safety
Auto-Remediation (Customer-Built)	Customer implements their own circuit-breakers using our API	Mature FinOps teams with existing automation

Auto-remediation capability can be built by customers on their infrastructure; we provide the detection signal, not the kill switch.

PART 2: THE TECHNICAL ROADMAP

Theoretical Limits & Performance Bounds

Bound Type	Current	Practical SOTA	Theoretical Max	Gap Closure Strategy
PR-AUC Ceiling	0.44	0.70–0.75	0.85–0.95	Full causal modeling (do-calculus on interventions)
Label Noise Impact	~20% noise	~5% noise	0% (oracle)	+0.10 via active learning + weak supervision
Unknown Unknowns	Supervised only	Semi-supervised	Fully unsupervised	+0.15 via GAN-based novelty detection
Temporal Resolution	Daily/Hourly	15-minute	Sub-second	+0.05 via streaming architecture

Implementation Priority (ROI-Based)

Priority	Technique	Effort	Expected Lift	ROI
1	LightGBM swap	1 day	+0.02–0.05	Very High
2	Cyclical + EMA features	2 days	+0.04–0.06	Very High
3	Pre-GNN graph features (HHI, Error Prop)	3 days	+0.03–0.05	High
4	SMOTE + Focal Loss	2 days	+0.04–0.06	High
5	Adaptive Baselines (Prophet/STL)	2 days	+0.04–0.08	High
6	Active Learning Loop	3 days	+0.05–0.08	High
7	Sub-hourly granularity	3 days	+0.03–0.05	Medium
8	LSTM per-app	1 week	+0.08–0.12	Medium
9	Ensemble + Conformal	3 days	+0.03–0.06	Medium
10	Change-Point Detection (BOCPD)	2 days	+0.05–0.07	Medium
11	Full GNN (GraphSAGE/GAT)	2 weeks	+0.05–0.10	Lower
12	Transformer (TFT/PatchTST)	3 weeks	+0.10–0.15	Lower

1. Real-Time Feasibility (Immediate Win)

Business Impact: Enables "early warning" alerts hours before cost spikes hit the bill.

Current State: Validated XGBoost Latency ($< 1 \mu\text{s}$ per sample). **The Limitation:** The model is currently run in batch mode (daily/hourly). **Proposed Solution:**

- **Stream Ingestion:** Deploy the trained XGBoost model to a Flink/Kafka consumer.
 - *Feasibility:* With inference latency at **0.80 microseconds**, the model can handle >1 million events per second on a single core, making it capable of processing the entire transaction firehose in real-time without sampling.

2. Advanced Handling of Class Imbalance

Business Impact: Reduces false negatives by 20-30%—catch more real anomalies before they escalate.

Current State: `scale_pos_weight` (Heuristic weighting). **The Limitation:** Simply weighting the minority class can lead to overfitting on noise. **Proposed Solution:**

- **SMOTE / ADASYN:** Use synthetic data generation to interpolate between existing anomalies. This creates a denser manifold for the decision boundary to learn from, rather than just point-weighting.
- **Focal Loss:** Implement the Focal Loss function (originally from RetinaNet for object detection [11]) to dynamically down-weight easy negatives and focus gradients on hard positives.
 - Equation: $FL(p_t) = -(1 - p_t)^\gamma \log(p_t)$

3. Frequency Domain Feature Engineering

Business Impact: Catches "first Monday of the month" patterns that cause 40% of false positives.

Current State: Time-domain lags ($t-1$, $t-24$). **The Limitation:** Lag features struggle to capture complex multi-seasonal patterns (e.g., "First Monday of the month"). **Proposed Solution:**

- **Fast Fourier Transform (FFT):** Decompose the time series into frequency components. Features representing the amplitude/phase of dominant frequencies (daily, weekly) allow the model to "subtract" the expected seasonality mathematically.
- **Wavelet Transforms:** For capturing transient anomalies (spikes) while preserving temporal localization.

3b. Cyclical Time Encoding & Calendar Features

The Gap: Tree models (XGBoost) treat `hour=23` and `hour=0` as distant values, when they are actually adjacent.

Proposed Solution:

- **Sin/Cos Encoding:** Transform cyclical features into continuous representations:

```
hour_sin = np.sin(2 * np.pi * hour / 24)
hour_cos = np.cos(2 * np.pi * hour / 24)
day_of_week_sin = np.sin(2 * np.pi * day_of_week / 7)
day_of_week_cos = np.cos(2 * np.pi * day_of_week / 7)
```

- **Holiday Flags:** Use `holidays` Python library to flag US/business holidays.
- **Pay Period Indicators:** First/last week of month often shows different spending patterns.
- *Expected Lift:* +0.02–0.04 PR-AUC

3c. Exponential Moving Averages (EMA) & Trend Detection

The Gap: Rolling means treat all past data equally; recent data should matter more. **Proposed Solution:**

- **Multi-horizon EMAs:** Compute `cost_ema_3d`, `cost_ema_14d`, `cost_ema_30d` per app.

```
pl.col('cost').ewm_mean(alpha=0.1).over('app_id').alias('cost_ema')
```

- **Trend Slope:** Linear regression slope of `value_per_cost` over 7-day window detects gradual degradation.
- **EMA Crossovers:** Flag when `cost_ema_3d > cost_ema_14d` (short-term acceleration signal).
- *Expected Lift:* +0.02–0.03 PR-AUC (catches "slow leak" anomalies before they become critical)

4. Topological Embeddings (Graph Neural Networks)

Business Impact: Predicts cascade failures—"Service B is failing, Service A will fail in 2 hours."

Current State: Aggregated stats (`unique_suppliers_count`). **The Limitation:** We lose the *structure* of the dependency graph. We know *how many* dependencies exist, but not *who* they are or their health state.

Proposed Solution:

- **GraphSAGE / GAT:** Train a Graph Neural Network to generate node embeddings.
 - *Why:* This allows a service to "inherit" the anomaly state of its neighbors. If Service A calls Service B, and Service B is failing, the GNN embedding for A will reflect this risk before A even shows symptoms.
- **Edge Weighting:** Use transaction volume and avg response time as edge features for richer graph structure.
- *Expected Lift:* +0.05–0.10 PR-AUC (requires PyG infrastructure)

4b. Pre-GNN Graph Features (Quick Wins)

The Gap: Full GNN implementation is complex; simpler graph metrics can provide immediate value.

Proposed Solution:

- **Herfindahl-Hirschman Index (Supplier Concentration):**

```
# HHI = Σ(supplier_cost_share²) – measures supplier diversity
# Low HHI = diversified suppliers (safer)
# High HHI = single-supplier dependency (risky)
supplier_shares = cost_per_supplier / total_cost
HHI = (supplier_shares ** 2).sum()
```

- **Error Propagation Score:** Percentage of suppliers with `error_rate > 0.1` feeding this app.
 - *Why:* Predicts cascade failures—if your suppliers are unhealthy, you will be too.
- **Lagged Network Metrics:** `unique_suppliers_7d_ago` to detect sudden dependency changes.
- **Supplier Failure Exposure:** Weighted average of supplier error rates by transaction volume.
- *Expected Lift:* +0.03–0.05 PR-AUC (graph structure currently underutilized)

5. Unsupervised Manifold Learning (Autoencoders)

Business Impact: Finds "unknown unknowns"—novel attack patterns or cost anomalies never seen before.

Current State: Supervised Learning (XGBoost) on weak labels. **The Limitation:** The model is limited by the quality of our heuristic labels. It cannot find "unknown unknowns." **Proposed Solution:**

- **LSTM-Autoencoder:** Train a neural network to compress and reconstruct the normal time series.

- *Mechanism:* The model learns to reconstruct "normal" patterns perfectly (low reconstruction error). When it sees a novel anomaly, it fails to reconstruct it (high reconstruction error).
- *Benefit:* This removes the dependence on "weak labels" entirely.
- *Expected Lift:* +0.08–0.12 PR-AUC

5b. Transformer-Based Sequence Models

The Gap: LSTMs struggle with very long sequences and lack interpretability. **Proposed Solution:**

- **Temporal Fusion Transformer (TFT):** [12]
 - Handles multiple time series with different characteristics simultaneously.
 - Built-in attention mechanism for interpretability ("why was this flagged?").
 - Designed for forecasting with known future inputs (holidays, scheduled events).
- **PatchTST:** [13]
 - State-of-the-art for time series anomaly detection (2023).
 - Patches time series like images, applies transformer attention.
 - Outperforms LSTMs on long-horizon tasks.
- *Expected Lift:* +0.10–0.15 PR-AUC (major improvement, but requires GPU infrastructure)
- *Trade-off:* More complex training, longer development cycle

6. Model Ensembling & Uncertainty Quantification

Business Impact: Reduces alert fatigue 30-40%—on-call engineers only see high-confidence alerts.

Current State: Single Point Estimate (XGBoost). **The Limitation:** We don't know how "confident" the model is. A probability of 0.6 could mean "ambiguous data" or "confident but weak signal." **Proposed Solution:**

- **Voting Ensemble:** Combine XGBoost (Gradient Boosting) with a Random Forest (Bagging) and a Neural Network. This reduces variance and stabilizes predictions.
 - **Conformal Prediction:** Instead of a single class label, output a *set* of labels with a guaranteed error rate (e.g., "With 95% confidence, this is Normal"). This is critical for "Human-in-the-loop" systems to reduce alert fatigue.
 - **Isolation Forest Ensemble:** Add unsupervised Isolation Forest scores as a feature or ensemble member.
 - *Why:* Robust to label noise—detects outliers without relying on potentially weak labels.
 - *Implementation:* Stack IF anomaly scores as input feature for XGBoost, or use voting.
 - *Expected Lift:* +0.03–0.08 PR-AUC
-

7. Algorithm Alternatives (Beyond XGBoost)

Business Impact: Same accuracy, 2x faster training—ship improvements weekly instead of monthly.

7a. LightGBM (Quick Win)

Current State: XGBoost as sole gradient boosting model. **The Gap:** LightGBM often matches or exceeds XGBoost with faster training. **Proposed Solution:**

- **Drop-in Replacement:** `lgb.LGBMClassifier(is_unbalance=True)`
- **Advantages:** Faster training, native categorical support, better hyperparameter response.

- *Expected Lift*: +0.02–0.05 PR-AUC

7b. TabNet (Attention-Based Tabular)

The Gap: Tree models don't automatically learn feature interactions. **Proposed Solution:**

- **TabNet:** Neural network with attention mechanism for tabular data.
- **Advantages:** Auto feature selection, built-in interpretability via attention weights.
- *Trade-off:* Slower inference (~10 μ s vs 1 μ s), but better on complex patterns.
- *Expected Lift*: +0.03–0.07 PR-AUC

7c. CatBoost

The Gap: `app_id` is treated as numeric or one-hot encoded. **Proposed Solution:**

- **CatBoost:** Native handling of categorical features with ordered target encoding.
- **Advantages:** Better handling of high-cardinality categoricals like `app_id`.
- *Expected Lift*: +0.02–0.04 PR-AUC

8. Sub-Hourly Granularity

Business Impact: Catches the "\$50K overnight GPU bill" 6 hours before it happens.

Current State: Data aggregated to daily/hourly windows. **The Limitation:** Hourly aggregation smooths out transient spikes and burst patterns. **Proposed Solution:**

- **15-minute Windows:** Re-aggregate transactions to 15-min buckets for high-frequency apps.

```
df.groupby([pl.col('app_id'),
            pl.col('timestamp').dt.truncate('15m')]).agg(...)
```

- **Burst Detection:** Count transactions in 5-minute sliding windows, flag $>3\sigma$ bursts.
- **Response Time Trends:** Track `avg_response_time` at 15-min granularity for latency spikes.
- **Consecutive Failure Chains:** Count sequential "failure" responses within 1-hour window.
 - *Why:* Three failures in a row is more concerning than three spread across 24 hours.
- *Implementation Priority:* Apply to top-10 apps by transaction volume first.
- *Expected Lift*: +0.03–0.05 PR-AUC (captures transient anomalies missed by daily data)

9. Adaptive Baselines & Dynamic Thresholds

Business Impact: Eliminates "Black Friday is not an anomaly" false positives automatically.

Current State: Static thresholds (e.g., z-score > 2.0). **The Limitation:** Business patterns evolve (new products, pricing changes, seasonality shifts). Static thresholds become stale. **Proposed Solution:**

- **Prophet/STL Decomposition:** Decompose time series into trend + seasonality + residual.

```
from prophet import Prophet
# Fit per-app baseline, predict next-day "expected" value
# Anomaly = |actual - predicted| / uncertainty_interval
```

- **Auto-Retraining on Feedback:** User marks alerts as true/false positive → model retrains nightly.
 - **Dynamic Threshold Adjustment:** Thresholds adapt based on recent volatility per app.
 - *Expected Lift:* +0.04–0.08 PR-AUC (AWS/GCP 2025 production systems use this)
-

10. Active Learning & Human-in-the-Loop

Business Impact: Model gets smarter weekly from analyst feedback—20-30% noise reduction in 6 weeks.

Current State: Model trained once on heuristic labels. **The Limitation:** No mechanism to incorporate analyst feedback or correct mistakes. **Proposed Solution:**

- **Active Learning Loop:**
 1. Model flags uncertain predictions ($0.4 < P < 0.6$).
 2. Analyst reviews and labels these edge cases.
 3. Model fine-tunes on validated labels weekly.
 - *Result:* 20-30% noise reduction over 4-6 weeks.
 - **Weak Supervision (Snorkel):** [15] Combine multiple noisy labeling functions (error rate, cost spike, outage correlation) into probabilistic labels.
 - *Expected Lift:* +0.05–0.08 PR-AUC (addresses label quality bottleneck)
-

11. Root-Cause Attribution & Explainability

Business Impact: Saves 1 engineer-week per month on false positive investigation.

Current State: Model outputs probability; no explanation. **The Limitation:** Analysts waste time investigating false positives without context. **Proposed Solution:**

- **SHAP Values:** Compute per-prediction feature contributions.

```
import shap
explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_test)
# Alert: "Anomaly driven by: cost_per_request +45%, error_rate +200%"
```

- **Differential Analysis:** Compare anomalous period to baseline period.
 - "Top changes: GPU node group +20%, egress to eu-west-1 +50%"
 - **Actionable Alerts:** Attach likely root cause to each alert.
 - *Expected Lift:* +0.03–0.05 in effective precision (reduces investigation time 40-60%)
-

12. Change-Point Detection (Structural Breaks)

Business Impact: Detects pricing changes and architecture shifts—"New DB tier, not an anomaly."

Current State: Assumes stationarity within rolling windows. **The Limitation:** Major events (new pricing, architecture changes) break baselines. **Proposed Solution:**

- **Bayesian Online Change-Point Detection (BOCPD):** [14]

```
# Detect when underlying data distribution changes
# Flag: "Baseline shift detected at 2024-09-15 for app SELENE"
```

- **Regime Switching Models:** Different parameters for "normal" vs "high-load" regimes.
 - **Integration:** When change-point detected, reset rolling baselines and alert as "regime shift" not "anomaly."
 - *Expected Lift:* +0.05–0.07 PR-AUC (reduces false positives from legitimate changes)
-

13. Auto-Remediation (Production Readiness)

Business Impact: Customer can enable auto-pause for dev clusters—direct cost savings, zero manual work.

Current State: Alert-only system. **The Limitation:** Anomalies detected but not acted upon automatically.

Proposed Solution:

- **Non-Production Auto-Actions:**
 - Idle dev/test clusters → auto-pause after 2-hour inactivity.
 - Runaway queries → auto-terminate after cost threshold.
- **Production Safe-Guards:**
 - Throttle (not terminate) non-critical workloads.
 - Alert + recommend action for critical services.
- **Rules Engine Integration:** Define cost policies that trigger automatic responses.

```
- rule: idle_cluster_shutdown
  condition: cluster.cost > $100/day AND cluster.cpu_utilization < 5%
  action: pause_cluster
  environment: [dev, staging]
```

- *Value:* Direct cost savings; AWS/GCP now offer this GA (2025).
-

14. Multi-Cloud Normalization

Business Impact: Unified visibility across AWS, Azure, and GCP in a single dashboard.

Current State: Single-tenant demo data. **The Limitation:** Real deployments span AWS/Azure/GCP with different billing lags and taxonomies. **Proposed Solution:**

- **Vendor Lag Reconciliation:** AWS bills 24-48h delayed; Azure near-real-time. Normalize to common timeline.

- **Tag Hygiene Enforcement:** Reject untagged resources from anomaly pipeline; enforce consistent naming.
- **Unified Cost Model:** Convert vendor-specific metrics to common schema:

```
# AWS: cost_and_usage_report → normalized_cost
# Azure: cost_management_export → normalized_cost
# GCP: billing_export → normalized_cost
```

- *Expected Lift:* +0.02–0.04 PR-AUC (cleaner data = better models)

15. Probabilistic Anomaly Scoring

Business Impact: Distinguishes "confident anomaly" from "uncertain, needs review"—smarter alert triage.

Current State: Point estimates ($P(\text{anomaly}) = 0.73$). **The Limitation:** No confidence intervals; can't distinguish "sure it's anomalous" from "uncertain." **Proposed Solution:**

- **Gaussian Processes:** Output mean + variance for each prediction.
 - High mean, low variance → confident anomaly.
 - High mean, high variance → uncertain, needs review.
- **Bayesian Neural Networks:** Monte Carlo dropout for uncertainty quantification.
- **Confidence-Based Alerting:**
 - High confidence → auto-alert.
 - Low confidence → queue for human review.
- *Expected Lift:* Reduces alert fatigue 30-40% by prioritizing confident predictions.

16. Phase 7: Dual-Brain Architecture (Long-Term Vision)

Business Impact: Fast alerts for critical spikes combined with deep analysis for subtle inefficiencies.

The Strategic Bottleneck: By Phase 6, deep learning models can achieve high accuracy but with heavier latency (~50–100ms). This may be too slow for real-time alerting on massive cost spikes. Conversely, fast tree models may miss subtle, slow-burning inefficiencies.

The Solution: Split the pipeline into two distinct systems—a pattern used by Waymo [16], OpenAI [17], Netflix [18], and Uber [19].

System 1: The Reflex (Fast Brain)


Attribute	Specification
Goal	"Detect the Large Tears" — alert on critical spikes immediately
Tech Stack	LightGBM + Polars (streaming via Kafka/Flink)
Latency	< 1 millisecond (inference time)
Target PR-AUC	~0.65 (optimized for speed)


Attribute	Specification
Optimization	Recall — don't miss disasters
Output	High-priority alerts; customers can optionally build auto-pause on their end
Industry Precedent	Waymo reactive control [16], Netflix fraud blocking [18]

System 2: The Cortex (Slow Brain)

Attribute	Specification
Goal	"Detect Microscopic Leaks" — find efficiency drains
Tech Stack	TFT [12] + Causal Graphs + Active Learning
Latency	Asynchronous (minutes/hours)
Target PR-AUC	0.85+ (with human-validated labels)
Optimization	Precision — don't waste engineer time on false positives
Output	Investigation reports, root cause analysis, architecture recommendations
Industry Precedent	OpenAI o1 reasoning [17], Uber Michelangelo [19]

Architecture Diagram





Industry Precedent for Dual-Brain Systems

Company	System 1 (Fast Reflex)	System 2 (Deep Reasoning)	Source
Waymo	Reactive control (<10ms braking/steering)	Transformer-based route planning	[16]
OpenAI	Fast token prediction (GPT-4)	Chain-of-thought reasoning (o1)	[17]
Netflix	Real-time fraud blocking (<ms latency via Kafka)	Behavioral pattern analysis (31% content fraud detection)	[18], [25], [26]
Uber	Real-time pricing/dispatch	Async demand forecasting (Michelangelo)	[19]

Netflix Case Study: Real-Time Fraud Detection

Netflix's fraud detection system provides a direct precedent for the Dual-Brain architecture:

System 1 (Reflex) Implementation:

- Real-time ML models on streaming events via **Kafka ingestion** [27]
- Detects login spikes, multi-location surges with **<ms latency**
- Uses **XGBoost + adaptive heuristics** for immediate blocking
- Handles **1M+ accounts** with semi-supervised models + **SMOTE** for imbalanced data [25], [26]

System 2 (Cortex) Implementation:

- Behavioral analytics identifies **31% of content fraud** (rapid license grabs) [18]
- Async pattern analysis for account sharing detection
- Human-in-the-loop validation for edge cases

This mirrors our LightGBM/Polars streaming (System 1) + TFT/Active Learning (System 2) architecture.

PART 3: VALIDATION & REFERENCES

17. Validation Protocol (For High PR-AUC Claims)

Before presenting any PR-AUC > 0.75, run these mandatory checks:

Required Validations

1. **Strict Time-Split:** Train on older data, test on newer (no shuffling)
2. **Leave-One-App-Out:** Ensure model generalizes to unseen apps
3. **Feature Importance Audit:** No single feature >30% importance
4. **Label Leak Check:** Verify no features derived from target or future data

Warning Signs of Data Leakage

- PR-AUC jumps from 0.60 to 0.95 with one new feature
 - `outage_cost` or `error_rate` has near-perfect importance
 - Model performs worse on held-out apps than seen apps
-

References

Cloud Cost & Anomaly Detection

1. [PMC: Machine Learning for Anomaly Detection](#)
2. [Ternary: Cloud Cost Anomaly Detection](#)
3. [FinOps: Managing Cloud Cost Anomalies](#)
4. [HPI: Anomaly Detection Survey \(Schmidl et al., 2022\)](#)
5. [arXiv: Recent Advances in Anomaly Detection](#)
6. [Google Cloud: Cost Anomaly Detection GA \(2025\)](#)
7. [AWS: Cost Anomaly Detection Updates \(2025\)](#)
8. [Binadox: AI-Powered Cloud Cost Monitoring 2025](#)
9. [CloudZero: Cloud Cost News May 2025](#)

Academic Foundations

10. [Kaggle IEEE-CIS Fraud Detection Competition](#)
11. [Lin et al.: Focal Loss for Dense Object Detection \(RetinaNet\)](#)
12. [Lim et al.: Temporal Fusion Transformers \(Google, 2021\)](#)
13. [Nie et al.: PatchTST - A Time Series is Worth 64 Words \(2023\)](#)
14. [Adams & MacKay: Bayesian Online Changepoint Detection](#)
15. [Snorkel AI: Programmatic Labeling & Weak Supervision](#)

Dual-Brain Architecture Precedents

16. [Waymo Safety Report: Autonomous Vehicle Architecture](#)
17. [OpenAI: Learning to Reason with LLMs \(o1\)](#)
18. [Netflix Tech Blog: Machine Learning for Fraud Detection in Streaming Services](#)
19. [Uber Engineering: Michelangelo ML Platform](#)

Netflix Fraud Detection (Detailed)

25. [InfoQ: Netflix's AI Fraud Detection Framework \(2022\)](#) — Semi-supervised models + SMOTE for imbalanced data
26. [GeeksforGeeks: How Netflix Uses Data Science](#)
27. [LinkedIn: How Apache Kafka Powers Netflix, Uber](#) — Real-time streaming architecture
28. [DataExpert: Scaling Netflix's Threat Detection](#)

Graph Neural Networks

- 20. [Pinterest: PinSage - Graph Convolutional Neural Networks \(KDD 2018\)](#)
- 21. [Alibaba: GNN for Fraud Detection in E-Commerce](#)
- 22. [LinkedIn: Graph Neural Networks for Recommendations](#)

Uncertainty Quantification

- 23. [Amazon Science: Conformal Prediction for Time Series](#)
- 24. [AstraZeneca: Conformal Prediction in Drug Discovery](#)

Cloud Cost Tool Comparisons

- 29. [CloudZero: Cloud Computing Statistics 2025](#)
- 30. [Scalr: Cloud Cost Optimization Best Practices 2025](#)